



# Scaling with

Frédéric Logé

[frederic.logemunerel@gmail.com](mailto:frederic.logemunerel@gmail.com)

2025-10-17




# Table of contents

- Introduction
- Writing efficient code
- Parallel treatment
- Analytics & analysis (part 1)
- Analytics & analysis on cluster
- Resources


# Introduction

# Data Science Lifecycle




Data Science Life cycle, [source](#)

# What are we going to work on today ?



Data Science Life cycle, [source](#)


# What are the tools involved ?



Data Science Life cycle, [source](#)

# Common issues with large datasets


Time, Memory, Crash



```
* Found 1372 raster cells that were NA for one or more, but not all, predictor
* Removed 94 occurrence localities that shared the same grid cell.
* Removed 2 occurrence points with NA predictor variable values.
* Removed 1 background points with NA predictor variable values.
* Assigning variable bhutanlulc to categorical ...
* Clamping predictor variable rasters...
* Model evaluations with random 10-fold cross validation...

*** Running ENMeval v2.0.3 with maxnet from maxnet package v0.1.4 ***

=====
r: cannot allocate vector of size 771.7 Mb
```



# But what can we do ?

A not exhaustive list of what we may do:

- **Optimize your R code** e.g. remove unnecessary loops, avoid data copy, loading unnecessary packages, etc, `Rcpp`
- **Rely on (most efficient) R packages** e.g. `dplyr`, `data.table`, `arrow`
- **Run code in parallel** e.g. `future` exploit hardware and distribute the work
- **Upgrade session's available memory** e.g. change RStudio config, get hardware update, setup virtual machine with higher memory config
- **Delegate treatment** e.g. `DBI`, `sparklyr`, `h2o` to perform operations with an engine more efficient than R
- **Work on data samples** e.g. active learning
- Breakdown tasks in your data science life cycle: **you cannot do everything in R.**

# Writing efficient code

# Best R practice recommendations

- Cleanliness & tidyness: avoid data copy, avoid garbage names throughout code, treat your RAM with kindness
- Comments are mandatory, even if variable names are explicit
- Work under RProject
- If you wish to build something that should last (used by others, robust etc), developing as an R Package is mandatory. Not suitable when doing tutorials and trying a million different things though :/

# Timing

When you are writing an R function, measuring it's execution time is good practice.

```
1 # define a function
2 foo <- function(){ return(sum(1:1e6)) }
3
4 # measure execution of the function, once
5 # user: actual CPU time for the process
6 # system: any indirect operation due to the process: I/O of files, GC, memory allocation, ...
7 # elapsed: total elapsed time
8 system.time({ foo() })
```

user	system	elapsed
0.000	0.000	0.001

```
1 # repeat 50 times the function to get some statistics
2 microbenchmark::microbenchmark({ foo() }, times = 50)
```


Unit: nanoseconds

	expr	min	lq	mean	median	uq	max	neval
{	foo()	164	164	13479.98	205	205	663831	50

# Profiling code

When your function (or general code) contains several instructions, profiling it allows you to see where exactly things go wrong. Here's a vanilla example:

```
1 # note that in RStudio you can use the "Profile" menu identically
2 profvis::profvis({
3   r <- c()
4   s <- 0
5   for(i in 1:1e6){
6     s <- s + i
7     r[i] <- i
8   }
9 })
```



<GC>

```
r[i] <- i  
doTryCatch  
tryCatchOne  
tryCatchList  
tryCatchList
```

Sample Interval: 10ms

130ms

# Exercise: Moving Average

## Moving Average

Input:  $x$ , numerical vector of length  $N$

Output:  $y$ , numerical vector of same length, where  $y_i := (x_{i-1} + x_i + x_{i+1})/3$  if  $i > 1$  and  $i < N$ , otherwise  $y_i$  is NA.

Propose some implementations of this function and use the timing functions seen previously to time yourselves.

# Solution(?)

```
1 x <- rnorm(n=1e6)
2 n <- length(x)
3
4 ma_0 <- function(x, n){
5   y <- rep(NA, n)
6   for(i in (2:(n-1))){
7     y[i] <- (x[i-1]+x[i]+x[i+1])/3
8   }
9   return(y)
10 }
11
12 ma_1 <- function(x, n){
13   return((c(NA, x[1:(n-1)]) + x + c(x[2:n], NA))/3)
14 }
15
16 microbenchmark::microbenchmark(ma_0(x=x, n=n), ma_1(x=x, n=n))
```

Unit: milliseconds

	expr	min	lq	mean	median	uq	max	neval
	ma_0(x = x, n = n)	70.60356	74.18450	77.02036	75.49027	77.47805	119.4600	100
	ma_1(x = x, n = n)	10.97320	13.21899	16.20814	14.16101	15.40983	45.8453	100

Vector operations are a must-use resource of R.

# Why still create our own R functions ?

If it exists in R and runs decently well, then why rebuild it ?

There are many cases were we might not find our pick:

- Data simulation purposes
  - Statistical metrics
  - Optimization functions
- ...
- Because we don't know any other language 😭

# Exercise: Simulation Example

## Simulation

Consider the following dynamic system:

- $x_0$  is in  $[0; 1]$
- $\lambda$  is in  $[0; 4]$

Then, for all  $t \geq 1$ :

$$x_t = \lambda * x_{t-1} * (1 - x_{t-1}).$$

Step 1. Code an R function which takes as input  $(x_0, \lambda, n)$  and returns output  $x_n$ .

Step 2. Run the function for a uniform grid of  $(x_0, \lambda)$ , the size of your choosing.

# Solution to step 1


```
1 library(Rcpp)
2
3 run_iteration <- function(n_iter, x0, lambda){
4   for(i in 1:n_iter){
5     x0 <- lambda * x0 * (1-x0)
6   }
7   return(x0)
8 }
9
10 cppFunction("
double run_iteration_cpp(int n_iter, double x0, double lambda) {
  for (int i = 0; i < n_iter; i++) {
    x0 = lambda * x0 * (1.0 - x0);
  }
  return x0;
}
")
13
14 microbenchmark::microbenchmark(run_iteration(n_iter=1e6, x=0.5, lambda=3.8), run_iteration_cpp(n_it
```

Unit: milliseconds

	expr	min	lq		
	mean	median	uq	max	neval
run_iteration(n_iter = 1e+06, x = 0.5, lambda = 3.8)	17.765874	19.007231			
run_iteration_cpp(n_iter = 1e+06, x = 0.5, lambda = 3.8)	2.634496	2.856409			
	19.441426	19.655154	19.833156	21.343288	100
	2.994166	3.042877	3.066574	3.728909	100

# Solution to step 2 - plot

See more info on this [over here](#).



# Memory

Aside from time benchmarks, scanning our environment for large objects can always be useful.

- `ls()` provides you the list of elements in your env
- `object.size(<the element>)` gives you its size (also `lobstr::obj_size(<the element>)`)
- `rm()` removes an object from the env
- `gc()` runs garbage collection (which runs periodically anyway so no need usually)

# Why do we need garbage collection ?

- RHS data is created and bounded to one or more names
- When a modification on RHS is requested, a copy is made.
- When an element is removed from the environment, it removes the name and its bind to the value, but the value in memory is still taken, until the garbage collector does its job.

```
1 a <- c(1, 2, 3)
2 b <- a # make copy
3
4 print(lobstr::obj_addr(a))
```

```
[1] "0x115b9cbe8"
```

```
1 print(lobstr::obj_addr(b))
```

```
[1] "0x115b9cbe8"
```

```
1 b[1] <- 0
2
3 print(lobstr::obj_addr(a))
```

```
[1] "0x115b9cbe8"
```

```
1 print(lobstr::obj_addr(b))
```

```
[1] "0x114694ce8"
```

# Exercise: Order of Magnitude

## Order of Magnitude

Generate random datasets:

- vary the data types: numerical, boolean, categorical
- vary the number of observations and columns

For each dataset generated, compute the object size and make a nice visualization from it

# Solution (?)

```
1 n <- 1e6  
2  
3 vec_1_million <- rnorm(n=n)  
4 print(lobstr::obj_size(vec_1_million))
```

8.00 MB

```
1 binary_vec_1_million <- round(rnorm(n=n))  
2 print(lobstr::obj_size(binary_vec_1_million))
```

8.00 MB

```
1 char_vec_1_million <- as.character(rnorm(n=n))  
2 print(lobstr::obj_size(char_vec_1_million))
```

8.00 MB

Conclusion ?

# To confuse you:

```
1 char_vec_1_million <- as.character(rnorm(n=n))
2 print(lobstr::obj_size(char_vec_1_million))
```

8.00 MB

```
1 toto <- list(char_vec_1_million, char_vec_1_million, char_vec_1_million)
2 print(lobstr::obj_size(toto))
```

8.00 MB

```
1 toto[[1]][1] <- 9
2 print(lobstr::obj_size(toto))
```

95.65 MB

```
1 banana <- "bananas bananas bananas"
2 print(lobstr::obj_size(banana))
```

136 B

```
1 print(lobstr::obj_size(rep(banana, 100)))
```

928 B

# Take-Home Exercises

- Write an efficient ifelse block statement taking as input a numerical vector, with the conditions of your choice
- Investigate confusing memory examples in slide above
- Rcpp implementation of optimization function e.g. Decision Tree

# Parallel treatment

# The `future` package

If you look for resources around parallel programming, you'll inevitably found the following names: `snow`, `parallel`, `foreach` and `future`. Those are not the only ones but clearly the most popular.

In all instances, the code looks somewhat (if not a little more complex) like the one we shall use from the `future` package:

```
1 library(future)
2 library(doFuture)
3
4 plan(multisession, workers=1) # of workers - to be changed
5 system.time({
6   x <- foreach(i = 1:4) %dofuture% {
7     Sys.sleep(2)
8   }
9 })
```

```
user  system elapsed
0.059  0.003  8.078
```

# All for one ?

```
1 plan(multisession, workers=4)
2
3 # choose a slow function
4 slow_fct <- function(x){ Sys.sleep(1e-5) ; log(x) }
5
6 # main vector
7 x <- c(1:1e5)
8
9 # iterate through the list
10 system.time({
11   iter <- 1:length(x)
12   foreach(i=iter, .combine='c') %dofuture% {
13     slow_fct(x[i])
14   }
15 })
```

user	system	elapsed
11.830	0.136	12.530

```
1 # iterate faster (especially in very large setting)
2 system.time({
3   iter <- itertools::isplitIndices(n=length(x), chunks=4)
4   foreach(i=iter, .combine='c') %dofuture% {
5     slow_fct(x[i])
6   }
7 })
```

user	system	elapsed
0.052	0.006	0.087

```
1 # baseline
2 system.time({
3   sapply(X=1:length(x), FUN=function(i){ slow_fct(x[i]) })
4 })
```

```
user  system elapsed
0.135    0.107   1.627
```

# Analytics & analysis (part 1)

# Data Wrangling

Recommended max data size	Package
~100K	<b>base R</b>
~1M	<b>readr, dplyr</b>
~10M	<b>data.table</b>
Inf	<b>arrow, duckdb</b>

# Analytics & analysis on cluster

db connection h2o spark

# Resources

# Resources

## Textbooks

Advanced R, 2nd Edition, utilitR book, Advanced R training

## Benchmarks

Data Wrangling Benchmark 1, Data Wrangling Benchmark 2

## Cheatsheets & documentation links

data.table, dplyr, sparklyr

## Miscellaneous

SAS to R